



Notes for classroom use

- 1) Baseline (base value) is the model's expected output on the background data; for tree models it is on the log-odds scale.
- 2) SHAP values ϕ represent signed feature contributions; $\text{base} + \sum \phi \approx \text{model output on the chosen scale}$.
- 3) Global views aggregate $|\phi|$ across instances for stable rankings; dependence plots reveal interactions.
- 4) Local views trace how large ϕ move an instance toward the final prediction (force/decision plots).
- 5) Use insights for feature monitoring, data quality checks, and model governance; SHAP is not causal evidence.